

Fine-Tuning Large Language Model with Custom Dataset for Ansible Code Generation

Furkan Gürbüz

Technical University of Munich

TUM School of Computation, Information and Technology

Chair of Information Systems and Business Process Management (i17)

Prof. Dr. Stefanie Rinderle-Ma

Thomas Teubner, Noah Kim, Dr. Holger Wittges

München, 13. Mai 2025



Increase efficiency and quality of IaC coding for DevOps

- Manual Ansible playbooks are complex and time-consuming
- Research lacks of domain-specific LLMs for the Academic Chair
- Sensitive information in the code sequences that are used as prompts can create issues with commercial LLM's



Ansible Playbooks are like recipe books for servers

Key concepts of Ansible

- Written in YAML code
- Defines tasks (e.g. install package xy)
- each playbook contains one or more “plays” (sets of tasks)
- each play targets specific machines and performs actions in a defined order.
- Target the servers where tasks should run

Simple Code snippet

```
- name: Say hello
hosts: localhost
tasks:
  - name: Print a message
    debug:
      msg: "May the Force be with you"
```

What methodology can be used to gather and prepare a custom dataset for fine-tuning large language models?

Methodology

- Created a custom Input-Output paired dataset
- Used specialized tools to gather Ansible code samples
- Cleaned and preprocessed data to remove noise
- Split into train, validation, and test sets for robust evaluation

Expected Result

- A clean, structured dataset ready for fine-tuning a large language model



What steps are involved in implementing the fine-tuning process for the large language model?

Methodology

- Selected a pre-trained LLM (Phi-4)
- Integrated LoRA for efficient fine-tuning
- Used Unsloth and PyTorch frameworks
- Executed training with the custom Ansible dataset

Expected Result

- A fine-tuned LLM optimized for Ansible code generation in SAP environments



To what extend does the fine-tuned LLM meet the requirements?

Methodology

- Evaluated with ROUGE, METEOR, CHRF, and Ansible-lint
- Measured accuracy, syntax correctness, and code quality

Expected Result

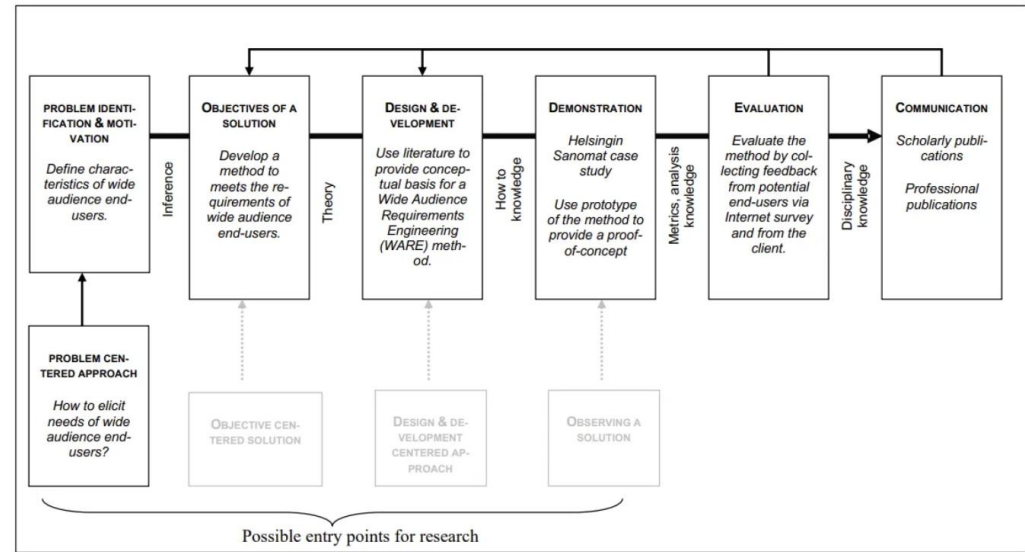
- Clear insights on how well the model performs in Ansible code generation tasks



Leaning on the DSR methodology of Hevner et al.

- Primary artifact is fine-tuned large language model
- Iterative cycles of evaluation and improvement in development

→ Artifact aims to enhance quality, efficiency and accuracy of code generation



Utilizing RedHat's Ansible-Content-Parser to create our custom dataset

- Extracted code data from existing GitHub repositories
- Split the dataset into a ratio of 70/15/15 (train/val/test)

```
{
  "data_source_description": "",
  "input": "
  ---\n# Ansible Playbook for SAP NetWeaver (JAVA) with IBM Db2 Sandbox
  installation\n\n# Use include_role / include_tasks inside Ansible Task block, instead of using
  roles declaration or Task block with import_roles.\n# This ensures Ansible Roles, and the tasks
  within, will be parsed in sequence instead of parsing at Playbook initialisation.\n\n\n####
  Begin Infrastructure-as-Code provisioning #####\n\n- name: Ansible Play to gather input for
  gathering vars and VM provisioning\n  hosts: localhost\n  gather_facts: false\n\n  # pre_tasks
  used only for Interactive Prompts only and can be removed without impact\n  pre_tasks:\n\n    -
  name: Playbook Interactive - Check if standard execution with an Ansible Extravars file is
  requested by end user",
  "license": "",
  "module": "ansible.builtin.set_fact",
  "output": "
    ansible.builtin.set_fact:\n
    playbook_enable_interactive_prompts: \n{{
  true if (sap_vm_provision_iac_type is undefined and sap_vm_provision_iac_platform is undefined)
  else false }}\n\n",
  "path": "deploy_scenarios/sap_nwas_java_ibmdb2_sandbox/ansible_playbook.yml",
  "repo_name": "ansible",
  "repo_url": "https://github.com/sap-linuxlab/ansible.playbooks_for_sap"
}
```



Microsoft's Phi-4 LLM is outperforming larger LLM's

- Phi-4 is a 14 billion parameter **transformer-based** LLM
- architecture of Phi-4 follows a **decoder-only** transformer

→ It reads the input prompt as part of the sequence, then it **predicts** one token at a time (also called **causal** or **auto-regressive**)

Benchmarks	Models						
	Phi-4 14B	Phi-3 14B	Qwen 2.5 14B instruct	GPT 4o-mini	LLaMA-3.3 70B instruct	Qwen 2.5 72B instruct	GPT 4o
MMLU	84.8	77.9	79.9	81.8	86.3	85.3	88.1
GPQA	56.1	31.2	42.9	40.9	49.0	50.6	50.6
MATH	80.4	44.6	75.6	73.0	66.3	74.6	74.6
HumanEval	82.6	67.8	72.1	86.2	78.9	87.1	90.6
MGSM	80.6	63.9	77.9	86.5	89.1	82.8	90.4
SimpleQA	3.0	7.6	7.6	39.4	9.3	8.6	9.3
DROP	75.5	58.3	59.7	79.9	82.4	80.9	85.6
MMLUPro	70.4	51.3	63.2	63.4	69.6	69.6	73.0
HumanEval+	82.8	69.2	79.1	82.4	77.8	84.0	88.0
ArenaHard	75.4	67.0	68.3	73.1	76.4	79.2	85.6
LiveBench	47.6	28.1	49.8	58.7	57.1	64.6	72.4
IFEval	63.0	57.9	78.7	78.7	89.3	85.6	84.8
PhiBench (internal)	56.2	43.9	49.8	58.7	57.1	64.6	72.4

Second Iteration metrics showcase significant increase in model performance

First Iteration Evaluation Scores	Second Iteration Evaluation Scores
'rouge1': np.float64(0.9030472630117133), 'rouge2': np.float64(0.8750916994262085), 'rougeL': np.float64(0.8937407874924173), 'rougeLsum': np.float64(0.8968060509310674)	'rouge1': np.float64(0.9089512020276276), 'rouge2': np.float64(0.8815313156096819), 'rougeL': np.float64(0.8997117333753821), 'rougeLsum': np.float64(0.9046659338280736)
'meteor': np.float64(0.8737622872894732)	'meteor': np.float64(0.8850545846138909)
'score': 97.2683986738109, 'char_order': 6, 'word_order': 0, 'beta': 2	'score': 97.84404601871951, 'char_order': 6, 'word_order': 0, 'beta': 2
Overall Ansible Lint Score: 0.56	Overall Ansible Lint Score: 0.76

Limitations

- Dataset inherits **code quality & security issues** from public GitHub sources
- No **semantic evaluation** by professional IaC developers
 - Evaluation lacks **real-world deployment testing**



Potential implementation and deployment of the fine tuned LLM on SAP BTP

- Embed the model in a **Fiori-based web application** on **SAP BTP**
- Deploy via **SAP AI Core** (side-by-side extensibility) or deploy via **FastAPI** or **Flask**, containerized with **Docker**
- Call the **REST API service** from the frontend using a **POST/GET request**



Thank you for your attention!

Questions?